

Bayesian_Inference_MNIST

November 9, 2025

Bayesian Inference on MNIST dataset

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from scipy.stats import multivariate_normal as mvn
```

Loading train and test files

```
[2]: #load train dataset
df=pd.read_csv('C:/Users/rishd/OneDrive/Desktop/Data Cleaning DS/MNIST/
                ↴MNIST_train.csv')
df
```

```
[2]:      Unnamed: 0    index    labels    0    1    2    3    4    5    6    ...   774   775   776   \
0            0        0        5    0    0    0    0    0    0    0    ...    0    0    0
1            1        1        0    0    0    0    0    0    0    0    ...    0    0    0
2            2        2        4    0    0    0    0    0    0    0    ...    0    0    0
3            3        3        1    0    0    0    0    0    0    0    ...    0    0    0
4            4        4        9    0    0    0    0    0    0    0    ...    0    0    0
...
59995      59995    59995        8    0    0    0    0    0    0    0    ...    0    0    0
59996      59996    59996        3    0    0    0    0    0    0    0    ...    0    0    0
59997      59997    59997        5    0    0    0    0    0    0    0    ...    0    0    0
59998      59998    59998        6    0    0    0    0    0    0    0    ...    0    0    0
59999      59999    59999        8    0    0    0    0    0    0    0    ...    0    0    0

    777   778   779   780   781   782   783
0            0    0    0    0    0    0    0
1            0    0    0    0    0    0    0
2            0    0    0    0    0    0    0
3            0    0    0    0    0    0    0
4            0    0    0    0    0    0    0
...
59995      0    0    0    0    0    0    0
59996      0    0    0    0    0    0    0
```

```
59997    0    0    0    0    0    0    0  
59998    0    0    0    0    0    0    0  
59999    0    0    0    0    0    0    0
```

[60000 rows x 787 columns]

```
[3]: #load test dataset  
df1=pd.read_csv('C:/Users/rishd/OneDrive/Desktop/Data Cleaning DS/MNIST/  
↪MNIST_test.csv')  
df1
```

```
[3]:      Unnamed: 0   index   labels   0   1   2   3   4   5   6   ...  774  775  776  777  \\\n0          0       0       7   0   0   0   0   0   0   0   ...   0   0   0   0   0  
1          1       1       2   0   0   0   0   0   0   0   ...   0   0   0   0   0  
2          2       2       1   0   0   0   0   0   0   0   ...   0   0   0   0   0  
3          3       3       0   0   0   0   0   0   0   0   ...   0   0   0   0   0  
4          4       4       4   0   0   0   0   0   0   0   ...   0   0   0   0   0  
...     ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...  
9995      9995    9995    2   0   0   0   0   0   0   ...   0   0   0   0   0  
9996      9996    9996    3   0   0   0   0   0   0   ...   0   0   0   0   0  
9997      9997    9997    4   0   0   0   0   0   0   ...   0   0   0   0   0  
9998      9998    9998    5   0   0   0   0   0   0   ...   0   0   0   0   0  
9999      9999    9999    6   0   0   0   0   0   0   ...   0   0   0   0   0  
  
    778   779   780   781   782   783  
0      0     0     0     0     0     0  
1      0     0     0     0     0     0  
2      0     0     0     0     0     0  
3      0     0     0     0     0     0  
4      0     0     0     0     0     0  
...   ...   ...   ...   ...   ...  
9995    0     0     0     0     0     0  
9996    0     0     0     0     0     0  
9997    0     0     0     0     0     0  
9998    0     0     0     0     0     0  
9999    0     0     0     0     0     0
```

[10000 rows x 787 columns]

```
[4]: df.shape
```

```
[4]: (60000, 787)
```

```
[5]: df1.shape
```

```
[5]: (10000, 787)
```

Dropping the first two index columns of train and test datasets

```
[6]: #dropping first two columns as they are index columns and will not be needed in
      ↪the analysis
df = df.iloc[:,2:]
df
```

	labels	0	1	2	3	4	5	6	7	8	...	774	775	776	777	778	779	\
0		5	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
1		0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
2		4	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
3		1	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
4		9	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
...
59995		8	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
59996		3	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
59997		5	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
59998		6	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
59999		8	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
		780	781	782	783													
0		0	0	0	0													
1		0	0	0	0													
2		0	0	0	0													
3		0	0	0	0													
4		0	0	0	0													
...													
59995		0	0	0	0													
59996		0	0	0	0													
59997		0	0	0	0													
59998		0	0	0	0													
59999		0	0	0	0													

[60000 rows x 785 columns]

```
[7]: #dropping first two columns of the test data
df1 = df1.iloc[:,2:]
df1
```

	labels	0	1	2	3	4	5	6	7	8	...	774	775	776	777	778	779	\
0		7	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
1		2	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
2		1	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
3		0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
4		4	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
...
9995		2	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
9996		3	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
9997		4	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0

```

9998      5  0  0  0  0  0  0  0  0  0  ...  0  0  0  0  0  0
9999      6  0  0  0  0  0  0  0  0  0  ...  0  0  0  0  0  0

    780  781  782  783
0      0  0  0  0
1      0  0  0  0
2      0  0  0  0
3      0  0  0  0
4      0  0  0  0
...
9995    0  0  0  0
9996    0  0  0  0
9997    0  0  0  0
9998    0  0  0  0
9999    0  0  0  0

```

[10000 rows x 785 columns]

Converting dataframe to a 2-D array for train and test datasets

```
[8]: #converting dataframe to an array for train data
X_train=df.to_numpy()
X_train
```

```
[8]: array([[5, 0, 0, ..., 0, 0, 0],
           [0, 0, 0, ..., 0, 0, 0],
           [4, 0, 0, ..., 0, 0, 0],
           ...,
           [5, 0, 0, ..., 0, 0, 0],
           [6, 0, 0, ..., 0, 0, 0],
           [8, 0, 0, ..., 0, 0, 0]])
```

```
[9]: #converting dataframe to an array for test data
X_test=df1.to_numpy()
X_test
```

```
[9]: array([[7, 0, 0, ..., 0, 0, 0],
           [2, 0, 0, ..., 0, 0, 0],
           [1, 0, 0, ..., 0, 0, 0],
           ...,
           [4, 0, 0, ..., 0, 0, 0],
           [5, 0, 0, ..., 0, 0, 0],
           [6, 0, 0, ..., 0, 0, 0]])
```

```
[10]: #Y is the target variable that is the feature label in this dataset
y_train=X_train[:,0]
y_train
```

```
[10]: array([5, 0, 4, ..., 5, 6, 8])
```

```
[11]: y_train.shape
```

```
[11]: (60000,)
```

```
[12]: #take all features excluding the first one for train dataset  
X_train = X_train[:,1:]  
X_train
```

```
[12]: array([[0, 0, 0, ..., 0, 0, 0],  
           [0, 0, 0, ..., 0, 0, 0],  
           [0, 0, 0, ..., 0, 0, 0],  
           ...,  
           [0, 0, 0, ..., 0, 0, 0],  
           [0, 0, 0, ..., 0, 0, 0],  
           [0, 0, 0, ..., 0, 0, 0]])
```

```
[13]: X_train.shape
```

```
[13]: (60000, 784)
```

```
[14]: #y for test data  
y_test = X_test[:,0]  
y_test
```

```
[14]: array([7, 2, 1, ..., 4, 5, 6])
```

```
[15]: y_test.shape
```

```
[15]: (10000,)
```

```
[16]: #take all features excluding the first one for test dataset  
X_test = X_test[:,1:]  
X_test
```

```
[16]: array([[0, 0, 0, ..., 0, 0, 0],  
           [0, 0, 0, ..., 0, 0, 0],  
           [0, 0, 0, ..., 0, 0, 0],  
           ...,  
           [0, 0, 0, ..., 0, 0, 0],  
           [0, 0, 0, ..., 0, 0, 0],  
           [0, 0, 0, ..., 0, 0, 0]])
```

```
[17]: X_test.shape
```

```
[17]: (10000, 784)
```

Scaling train and test data

```
[18]: #scale the training and test dataset to improve accuracy  
#we divide each value by 255 as there are 255 pixels  
X_train_scaled = X_train.astype(np.float32) / 255.0  
X_test_scaled = X_test.astype(np.float32) / 255.0
```

Visualizing MNIST Dataset

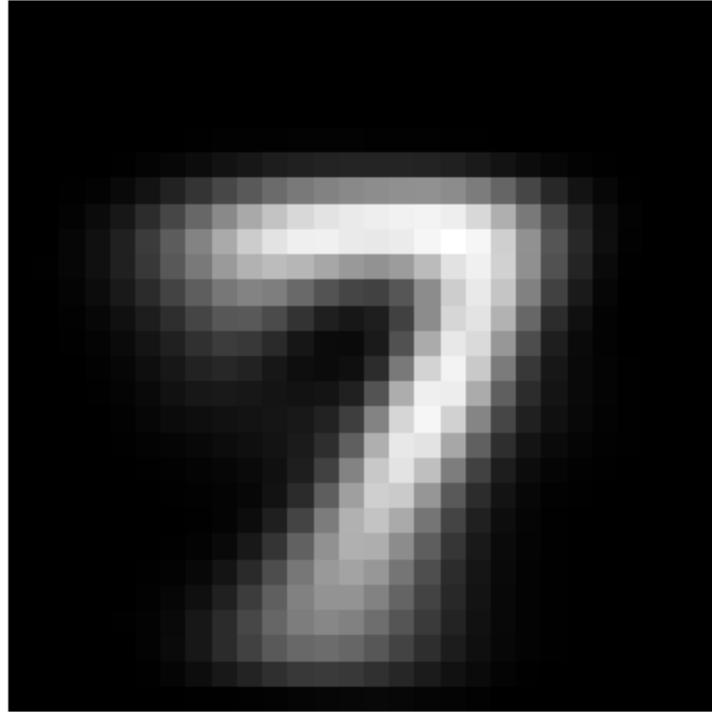
```
[19]: #img = X_train_scaled[0]           # takes first row  
#label = y_train[0]                 #takes first digit  
  
#img = img.reshape(28, 28)          #converts single 784 vector to 28*28 matrix  
  
#plt.imshow(img, cmap='gray')  
#plt.title(f"Label: {label}")  
#plt.axis('off')  
#plt.show()  
  
def show_me(X):  
    plt.imshow(X.reshape(28,28),cmap='gray')  
    plt.axis('off')  
    plt.show()  
  
show_me(X_train_scaled[135])
```



```
[20]: #average of all sevens in the dataset
```

```
def show_me_all_mean(X,y,k):
    show_me(sum(X[y==k,:]/len(X[y==k,:])))

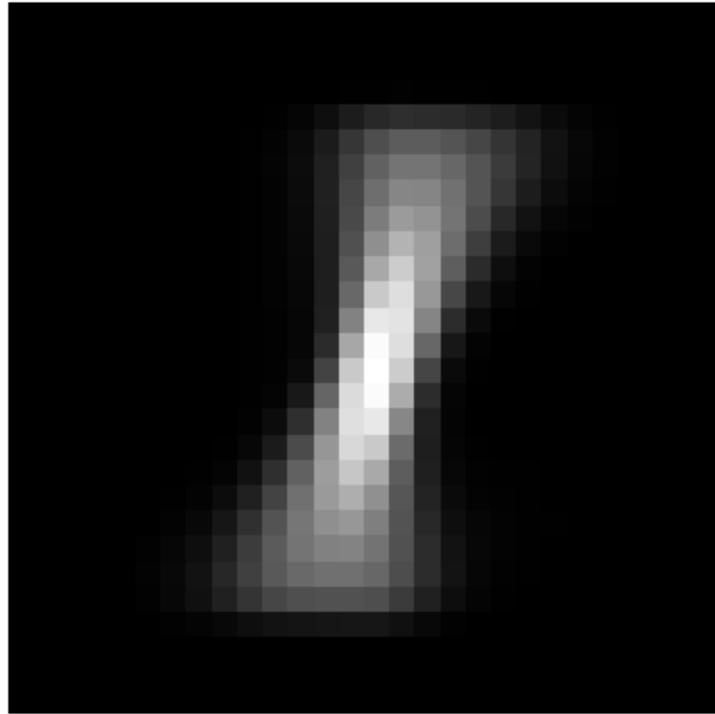
show_me_all_mean(X_train_scaled,y_train,7)
```



```
[21]: #average of all ones in the dataset
```

```
def show_me_all_mean(X,y,k):
    show_me(sum(X[y==k,:]/len(X[y==k,:])))

show_me_all_mean(X_train_scaled,y_train,1)
```



Naives Bayes Classifier

```
[22]: class GaussNB():      #contains methods and variables
    def fit(self,X,y,epsilon=1e-1):                      #hyper paramter tuned
        ↪to 1e-1, added to variance to prevent division by zero
        self.likelihoods=dict()                            #stores mean and var
        self.priors=dict()                                #stores prior probabilities

        self.K = set(y.astype(int))                      #stores labels as unique digits
        ↪from 0-9

        for k in self.K:
            X_k = X[y==k]
            self.likelihoods[k] = {"mean":X_k.mean(axis=0),"cov":X_k.
            ↪var(axis=0)+epsilon}           #P(X/y)
            self.priors[k] = len(X_k)/len(X)                #P(y)

    def predict(self,X):
        N,D = X.shape
        p_hat = np.zeros((N,len(self.K)))               #needs to be a tuple
```

```

    for k,l in self.likelihoods.items():
        p_hat[:,k] = mvn.logpdf(X,l["mean"],l["cov"]) + np.log(self.
        priors[k]) #log (X/y)*P(y) = log P(X/y)+log P(y)

    return p_hat.argmax(axis=1) #returns index of highest
    ↵posterior probability

```

[23]: def accuracy(y,y_hat):
 return np.mean(y==y_hat)

[24]: gnb = GaussNB()

[25]: #fit the model on train data
gnb.fit(X_train_scaled,y_train)

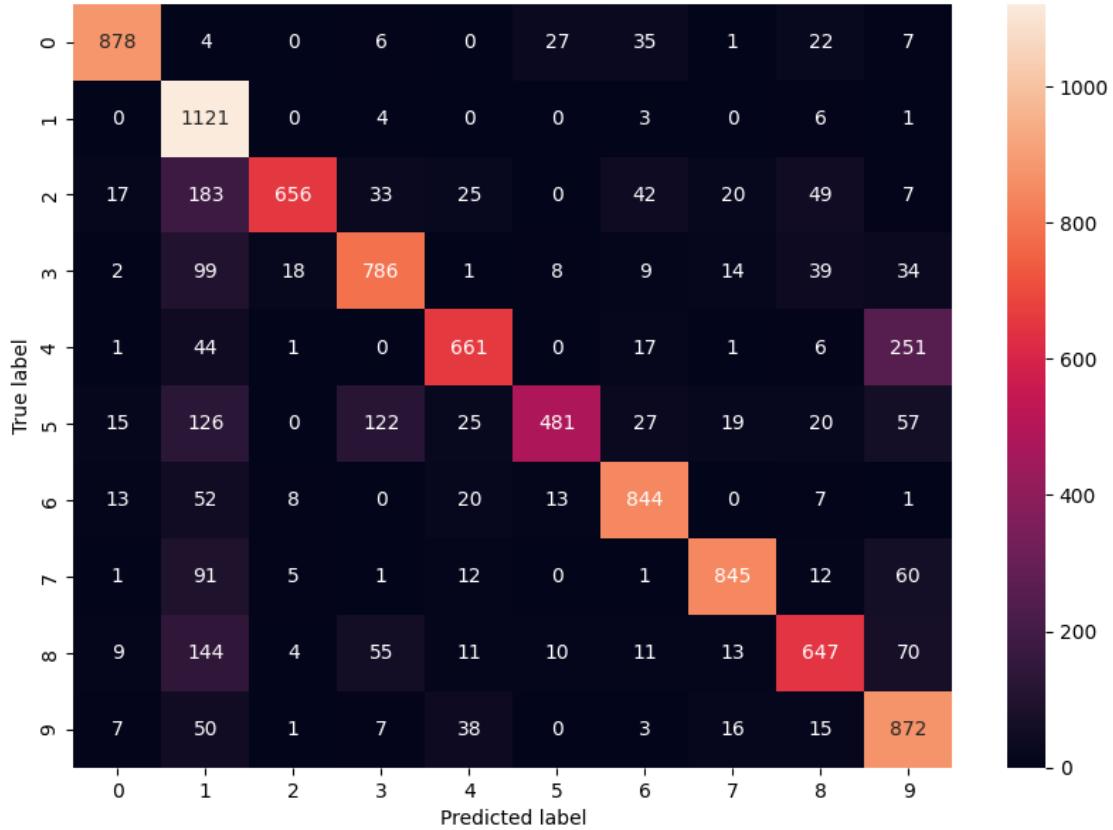
[26]: #make predictions on the test data
y_hat = gnb.predict(X_test_scaled)

[27]: #compare accuracies of predicted and actual test data labels
accuracy(y_test,y_hat)

[27]: np.float64(0.7791)

[28]: #confusion matrix for model evaluation
def ConfusionMatrix(y_actual,y_predicted):
 plt.figure(figsize=(10,7))
 y_actu = pd.Series(y_actual, name='Actual')
 y_pred = pd.Series(y_predicted, name='Predicted')
 cm = pd.crosstab(y_actu, y_pred) #creates a pivot table
 ↵that tells us how y_yest and y_hat relate with each other
 ax = sns.heatmap(cm, annot=True, fmt="d")
 plt.ylabel('True label')
 plt.xlabel('Predicted label')

ConfusionMatrix(y_test,y_hat)



Non-Naive Bayes Classifier

```
[29]: class GaussBayes():

    def fit(self, X, y, epsilon=1e-3):                      #fit method

        self.likelihoods = dict()                                #likelihood probability
        self.priors = dict()                                    #prior probability

        self.K = set(y.astype(int))

        for k in self.K:
            X_k = X[y==k,:]
            N_k, D = X_k.shape
            mu_k = X_k.mean(axis=0)

            self.likelihoods[k] = {"mean": mu_k,
                                   "cov": (1/(N_k-1)) * np.matmul((X_k-mu_k).T, np
                                   ↪X_k-mu_k)+epsilon*np.identity(D)}           #compute covariance
            self.priors[k] = len(X_k)/len(X)
```

```

def predict(self, X):                      #predict method
    N, D = X.shape
    P_hat = np.zeros((N, len(self.K)))

    for k,l in self.likelihoods.items():
        P_hat[:,k] = mvn.logpdf(X, l["mean"],l["cov"])+np.log(self.priors[k])
        ↵                                #posterior probability

    return P_hat.argmax(axis=1)

```

[30]: *#fit the model on train dataset*
gbays = GaussBayes()
gbays.fit(X_train_scaled,y_train)

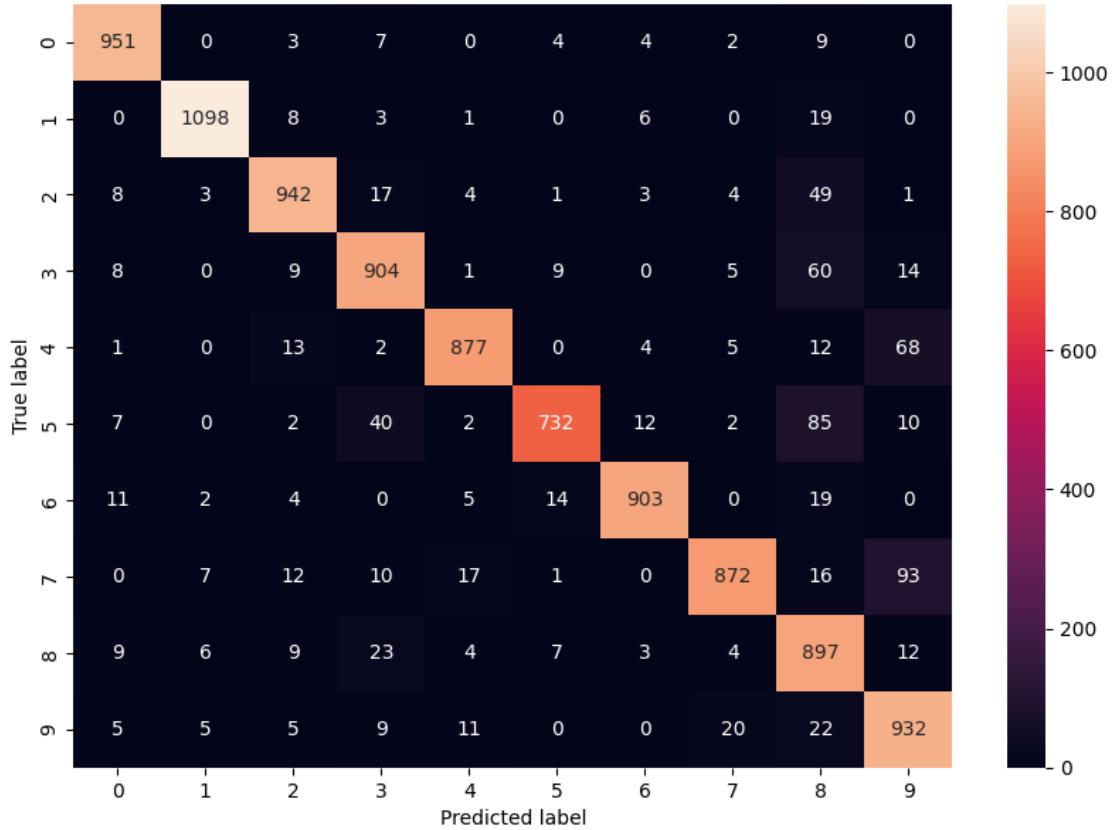
[31]: *#predict the model on the test dataset*
y_hat_bayes = gbays.predict(X_test_scaled)

[32]: accuracy(y_test,y_hat_bayes)

[32]: np.float64(0.9108)

[33]: *#confusion matrix for model evaluation*
def ConfusionMatrix(y_actual,y_predicted):
 plt.figure(figsize=(10,7))
 y_actu = pd.Series(y_actual, name='Actual')
 y_pred = pd.Series(y_predicted, name='Predicted')
 cm = pd.crosstab(y_actu, y_pred)
 ax = sns.heatmap(cm, annot=True, fmt="d")
 plt.ylabel('True label')
 plt.xlabel('Predicted label')

ConfusionMatrix(y_test,y_hat_bayes)



KNN Classifier

```
[ ]: #not included in the ppt as non-naive bayes model performed well on MNIST
    ↵dataset
```

```
[34]: class KNNClassifier():

    def fit(self,X,y):
        self.X = X
        self.y = y

    def predict(self,X,K,epsilon=1e-3):
        N = len(X)
        y_hat = np.zeros(N)

        for i in range(N):
            dist2 = np.sum((self.X-X[i])**2, axis=1)           #not taking sqrt cz it
            ↵is a slow process
            idxt = np.argsort(dist2)[:K]                      #returns sorted indices of
            ↵nearest neighbors
```

```

gamma_k = 1/(np.sqrt(dist2[idxt]+epsilon))

y_hat[i] = np.bincount(self.y[idxt],weights=gamma_k).argmax()

return y_hat

```

[35]: knn_instance = KNNClassifier()
knn_instance.fit(X_train_scaled,y_train)

[36]: y_hat_knn = knn_instance.predict(X_test_scaled,K=20)

[37]: accuracy(y_test,y_hat_knn)

[37]: np.float64(0.9633)

[38]: #confusion matrix
def ConfusionMatrix(y_actual,y_predicted):
 plt.figure(figsize=(10,7))
 y_actu = pd.Series(y_actual, name='Actual')
 y_pred = pd.Series(y_predicted, name='Predicted')
 cm = pd.crosstab(y_actu, y_pred)
 ax = sns.heatmap(cm, annot=True, fmt="d")
 plt.ylabel('True label')
 plt.xlabel('Predicted label')

ConfusionMatrix(y_test,y_hat_knn)

